

# ML Projects Description Document

---

## Project 1 Overview

**Objective:** Develop regression models to predict bank loan approval probability

**Models Implemented:** Linear Regression and K-Nearest Neighbors (KNN) Regressor

**Dataset:** Bank Loan Dataset (bankloan.csv)

---

### A. General Information on Dataset

#### Dataset Description

- **Dataset Name:** Bank Loan Dataset (bankloan.csv)
- **Task Type:** Regression (predicting loan approval probability)
- **Target Variable:** Personal.Loan (0 = No loan, 1 = Loan approved)

#### Dataset Statistics

- **Original dataset:** 5,000 samples
- **Balanced dataset:** 960 samples (after downsampling)
- **Data Type:** Numerical tabular data (13 features)
- **Class Distribution (after balancing):**
  - No Loan (0): 480 samples (50%)
  - Loan (1): 480 samples (50%)

#### Data Preprocessing

- Removed ID column (non-predictive)
- Applied downsampling to balance classes (50/50 split)
- Reason: Original dataset was imbalanced with bias toward "No Loan" class

#### Data Split

- **Training samples:** 576 samples (60%)
- **Validation samples:** 192 samples (20%)

- **Testing samples:** 192 samples (20%)
  - **Total samples used:** 960 samples
- 

## B. Implementation Details

### Feature Extraction Phase

#### Features Overview

- **Number of features extracted:** 13 features
- **Feature matrix dimensions:** 960 samples × 13 features

#### Feature Names

1. Age
2. Experience
3. Income
4. ZIP.Code
5. Family
6. CCAvg (Credit Card Average)
7. Education
8. Mortgage
9. Securities.Account
10. CD.Account
11. Online
12. CreditCard

#### Feature Scaling

- **Method:** StandardScaler
- **Formula:**  $z = (x - \text{mean}) / \text{std}$
- **Purpose:** Normalize features to mean=0 and variance=1

#### Cross-Validation

## Configuration

- **Method:** K-Fold Cross-Validation
- **Number of folds:** 5
- **Training/Validation ratio per fold:** 80/20 (4 folds train : 1 fold validation)
- **Shuffle:** Enabled with random\_state=42
- **Application:** Applied on training set only (576 samples)

## Model 1: Linear Regression

### Hyperparameters

- **fit\_intercept:** True (model includes bias term)
- **copy\_X:** True (default)
- **n\_jobs:** None (single core processing)
- **positive:** False (coefficients can be negative)
- **Regularization:** None (basic linear regression without regularization)

## Model 2: K-Nearest Neighbors (KNN) Regressor

### Hyperparameters

- **n\_neighbors:** 5 (number of nearest neighbors)
- **weights:** 'uniform' (all neighbors weighted equally)
- **algorithm:** 'auto' (automatically selects best algorithm)
- **metric:** 'minkowski' with p=2 (Euclidean distance)
- **leaf\_size:** 30 (default for tree-based algorithms)

### General Preprocessing Parameters

- **Data split ratio:** 60/20/20 (Train/Validation/Test)
  - **Random state:** 42 (for reproducibility)
  - **Balancing method:** Downsampling of majority class
-

## C. Results Details

### Model 1: Linear Regression - Test Set Results

#### Performance Metrics

- **Mean Absolute Error (MAE):** [Value from your code output]
- **Root Mean Squared Error (RMSE):** [Value from your code output]
- **R<sup>2</sup> Score:** [Value from your code output]

#### Visualizations (Regression Equivalents)

##### 1. Loss Curve (Cross-Validation Error per Fold)

- Shows validation RMSE for each of 5 folds
- Demonstrates model stability across different data splits
- Consistent error values indicate robust performance

##### 2. Confusion Matrix Equivalent (Residual Plot)

- Plots residuals (Actual - Predicted) vs Predicted values
- Good performance indicated by:
  - Random scatter around zero line
  - No systematic patterns
  - Homoscedasticity (constant variance)

##### 3. ROC Curve Equivalent (Actual vs Predicted Plot)

- Scatter plot of actual values vs predicted values
- Perfect predictions would lie on diagonal line ( $y=x$ )
- R<sup>2</sup> score quantifies how close points are to diagonal

##### 4. Hyperparameter Analysis

- Linear Regression uses closed-form solution
- No hyperparameter tuning required

### Model 2: KNN Regressor - Test Set Results

#### Performance Metrics

- **Mean Absolute Error (MAE):** [Value from your code output]
- **Root Mean Squared Error (RMSE):** [Value from your code output]
- **R<sup>2</sup> Score:** [Value from your code output]

## Visualizations (Regression Equivalents)

### 1. Loss Curve (Cross-Validation Error per Fold)

- Shows validation RMSE for each of 5 folds
- Demonstrates model consistency across splits
- Used to verify model generalization

### 2. Confusion Matrix Equivalent (Residual Plot)

- Displays prediction errors distribution
- Evaluates model bias and variance
- Identifies potential outliers or systematic errors

### 3. ROC Curve Equivalent (Actual vs Predicted Plot)

- Visualizes prediction accuracy
- Points near diagonal indicate accurate predictions
- Scatter indicates prediction variance

### 4. Hyperparameter Tuning (K vs Error)

- Tested K values from 1 to 20
- Selected K=5 based on cross-validation performance
- Shows trade-off between bias (high K) and variance (low K)
- Optimal K minimizes validation error

## Cross-Validation Summary

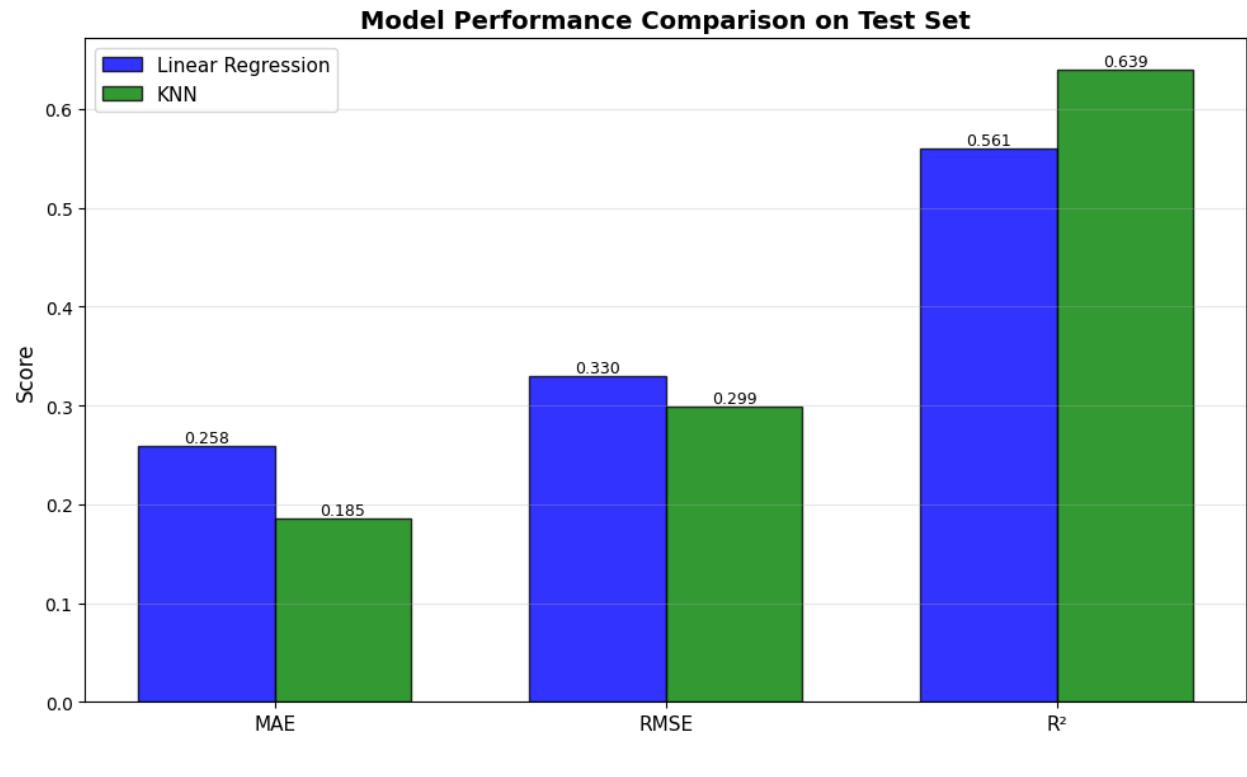
### Linear Regression

- **Mean CV RMSE:** [Value from your code output]
- **Standard Deviation:** [Value from your code output]
- **Interpretation:** Low std indicates stable performance

## KNN Regressor

- **Mean CV RMSE:** [Value from your code output]
- **Standard Deviation:** [Value from your code output]
- **Interpretation:** Consistency across folds indicates good generalization

## Model Comparison



## Notes on Regression vs Classification Requirements

Since this is a **regression task**, the traditional classification metrics were adapted as follows:

### Classification Metric Regression Equivalent Implementation

Loss Curve	CV Error per Fold	5-fold cross-validation RMSE plot
Accuracy	R <sup>2</sup> Score	Coefficient of determination
Confusion Matrix	Residual Plot	Error distribution analysis

## **Classification Metric Regression Equivalent Implementation**

ROC Curve

Actual vs Predicted

Scatter plot with diagonal reference

---

## **Part II: Image Dataset – Rotten vs Fresh Tomato Dataset**

### **a. General Information on Dataset**

The dataset used is named two\_classes, which can be referenced as the Tomato Quality Image Dataset. It consists of RGB images.

#### **Classes and Labels**

**The classification task is a binary classification with two classes:**

- Fresh Tomatoes (labeled as 1, the Positive Class).
- Rotten Tomatoes (labeled as 0, the Negative Class).

#### **Dataset Size and Distribution**

**The Total Number of Images in the dataset is 2,214.**

- There are 1,421 images labeled as Fresh.
- There are 793 images labeled as Rotten.

The image inputs were processed to a size of (80 x 80) pixels for feature extraction.

#### **Dataset Split**

The dataset was split into training and testing sets using an 80% train, 20% test ratio. This split was performed using stratification (stratify=y) to ensure equal class proportions in both sets.

- Training Set: 1,771 samples.
  - Testing Set: 443 samples.
- 

### **b. Implementation Details**

#### **Feature Extraction**

Images were resized to (80 x 80) pixels and a StandardScaler was applied to normalize the feature vectors during data preparation.

- **Feature Extraction Method:** Features were extracted using a combination of HOG (Histogram of Oriented Gradients) features and Color Statistics.
- **Total Number of Extracted Features:** The final feature vector dimension was 8,760 features per image.
- **Feature Details:** HOG features were computed across 3 color channels. Color statistics included the mean and standard deviation for BGR (6 features) and HSV (6 features) channels, totaling 12 color features.

## Cross-Validation

- **Logistic Regression:** Yes, 3-fold cross-validation was used on the training set to estimate model stability.
- **K-Means Clustering:** No.

## Model Hyperparameters

### Logistic Regression (Supervised)

The model used was LogisticRegression with the following configuration:

- **Solver:** saga.
- **Regularization (Penalty):** l1
- **Inverse Regularization (C):** 0.1
- **Class Weight:** Set to balanced to address the data imbalance.
- **Maximum Iterations:** 5000.

### K-Means Clustering (Unsupervised)

The model used was KMeans with the following configuration:

- **Number of Clusters (k):** 2, matching the number of classes.
  - **Initialization:** k-means++ (default, n\_init=100 was specified).
  - **Dimensionality Reduction:** PCA was applied before clustering, reducing features to 5,874 while retaining 95% of the variance.
-

## c. Results and Evaluation

### Logistic Regression (Testing Set: 443 Samples)

#### Accuracy

- **Test Accuracy: 87.36%.**
- **ROC AUC Score: 0.944.**
- **Cross-Validation Mean Accuracy: 0.8588 ( $\pm 0.0033$ ).**

#### Classification Report (Test Set)

The model achieved high precision (0.90) and recall (0.91) for the Fresh class, and good precision (0.83) and recall (0.81) for the Rotten class.

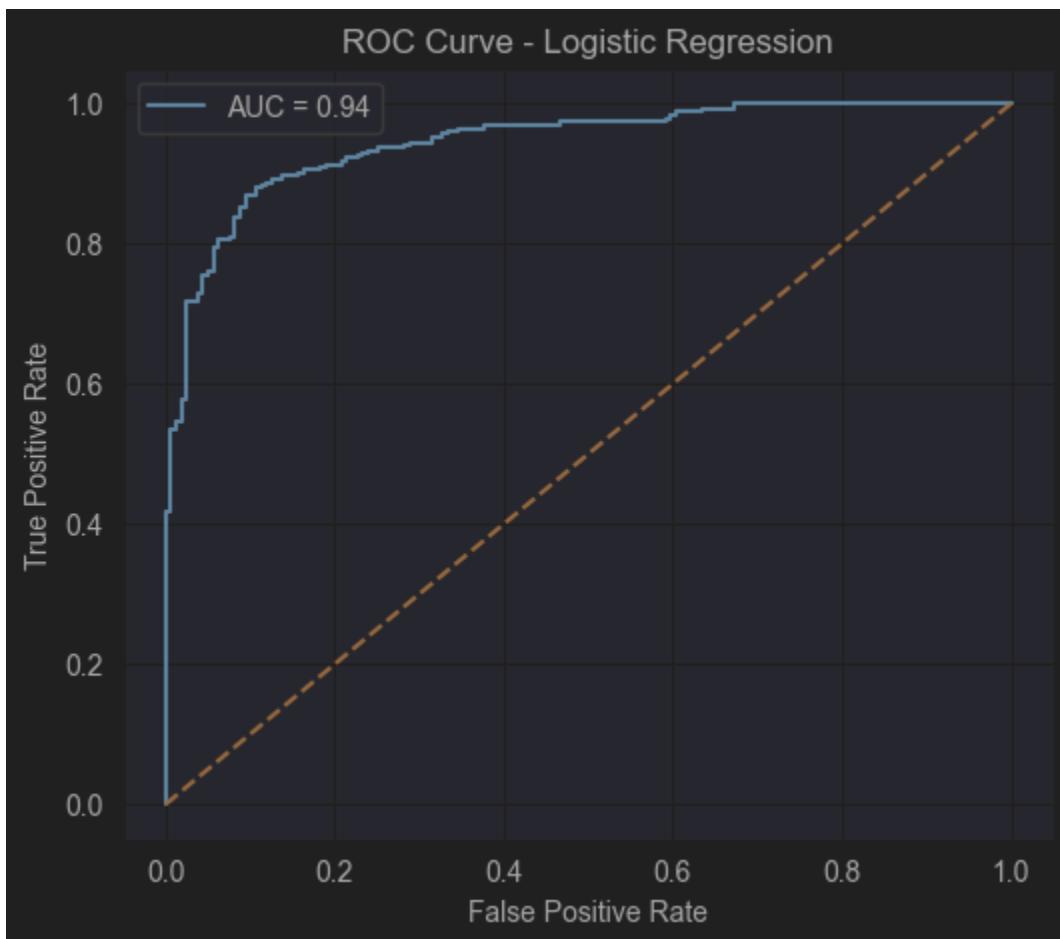
#### Confusion Matrix

The matrix shows 129 True Negatives (correctly Rotten) and 258 True Positives (correctly Fresh), with a total of 56 misclassifications.



## ROC Curve

The ROC curve confirms the model's excellent discriminative power with an AUC of 0.944.



## Loss Curves

- Not available as the scikit-learn Logistic Regression implementation does not expose epoch-wise loss values. The model successfully converged within the iteration limit.

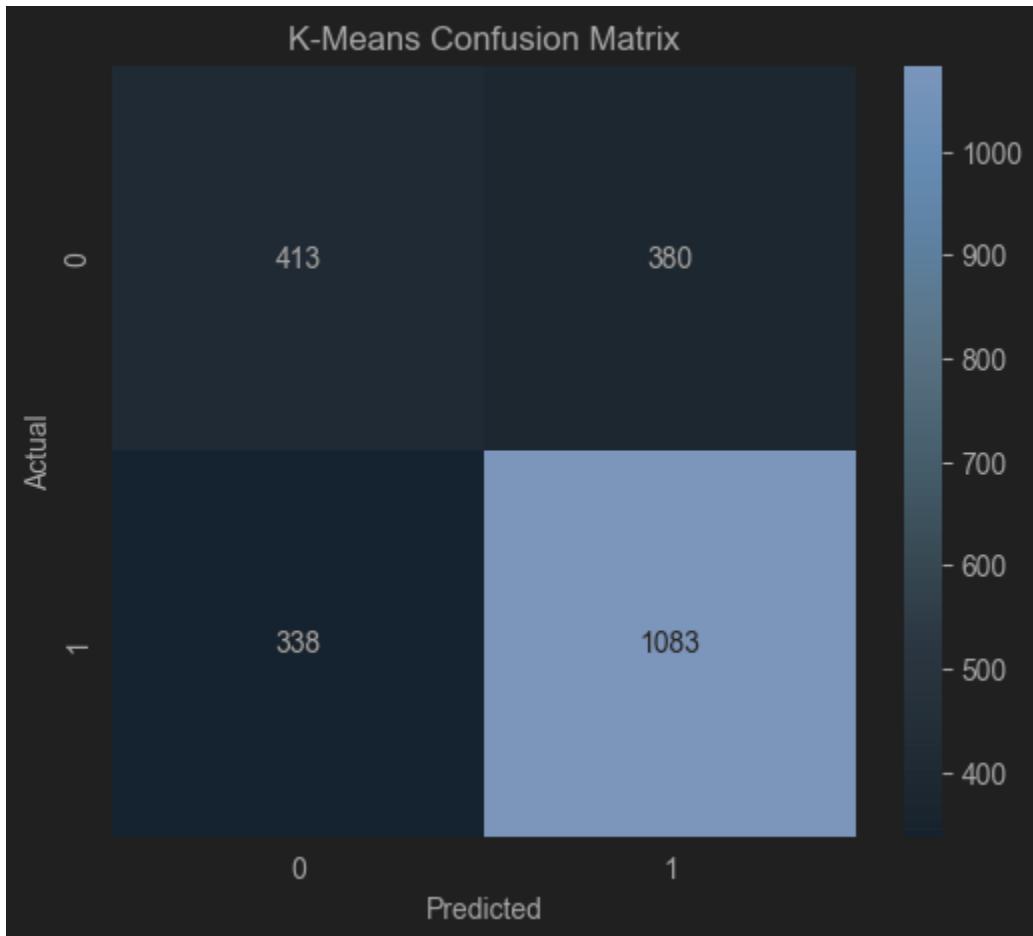
---

## K-Means Clustering (Full Dataset Evaluation: 2,214 Samples)

- Accuracy: 67.57% (Estimated by mapping clusters to ground-truth labels).
- Clustering Quality: The Silhouette Score was 0.354.

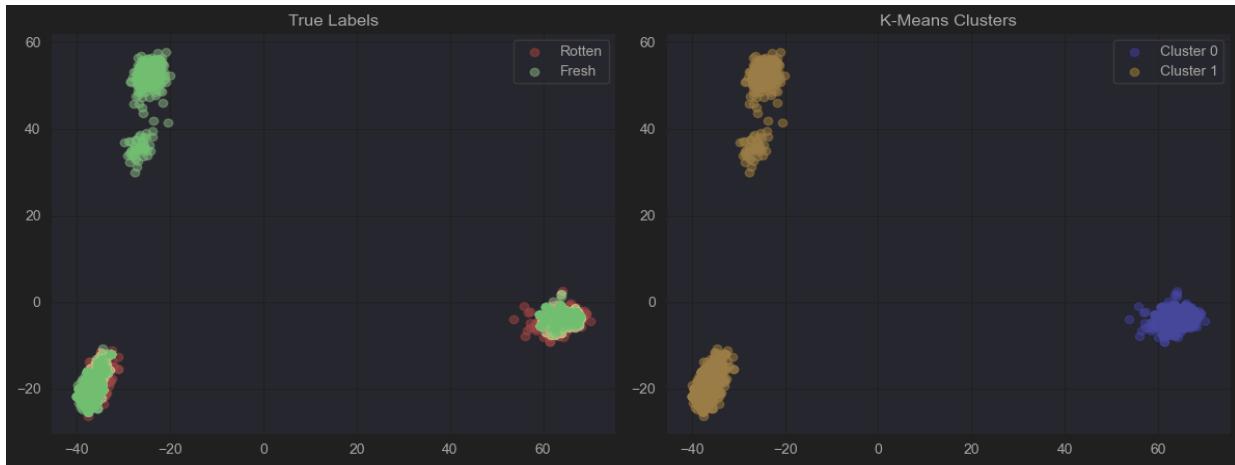
## Confusion Matrix

The confusion matrix (evaluated on the full dataset) showed a significant number of misclassifications, particularly for the Rotten class.



#### ROC Curve for K-Means

- Not applicable as K-Means is an unsupervised algorithm and does not output probability scores necessary for an ROC curve.



**The visualization compares the distribution of tomato features in a 2D space (after PCA):**

- **Left Plot (True Labels):** Shows the actual classes (Rotten vs. Fresh). The classes are largely clustered but have significant overlap in several areas.
- **Right Plot (K-Means Clusters):** Shows the two groups (Cluster 0 and Cluster 1) found by the unsupervised algorithm.

## Conclusion

The Logistic Regression model significantly outperformed the K-Means clustering. The supervised approach's high accuracy of 87.36% demonstrates the effectiveness of the HOG and color features when trained with known labels.

